

Multiplayer Reach-Avoid Games via Pairwise Outcomes

Mo Chen, Zhengyuan Zhou and Claire J. Tomlin

Abstract—A multiplayer reach-avoid game is a differential game between an attacking team with N_A attackers and a defending team with N_D defenders playing on a compact domain with obstacles. The attacking team aims to send M of the N_A attackers to some target location, while the defending team aims to prevent that by capturing attackers or indefinitely delaying attackers from reaching the target. The analysis of this game plays an important role in many applications. The optimal solution to this game is computationally intractable when $N_A > 1$ or $N_D > 1$. In this paper, we present two approaches for the $N_A = N_D = 1$ case to determine pairwise outcomes, and a graph theoretic maximum matching approach to merge these pairwise outcomes for an $N_A, N_D > 1$ solution that provides guarantees on the performance of the defending team. We will show that the four-dimensional Hamilton-Jacobi-Isaacs approach allows for real-time updates to the maximum matching, and that the two-dimensional “path defense” approach is considerably more scalable with the number of players while maintaining defender performance guarantees.

I. INTRODUCTION

Multiplayer reach-avoid games are differential games between two adversarial teams of cooperative players playing on a compact domain with obstacles. The “attacking team” aims to send as many team members, called “attackers”, to some target set as quickly as possible. The “defending team” seeks to delay or prevent the attacking team from doing so by attempting to capture the attackers. Such differential games have been studied extensively [1], [2] and are also powerful theoretical tools for analyzing realistic situations in robotics, aircraft control, security, and other domains [3], [4], [5].

The multiplayer reach-avoid game is difficult to analyze because the two teams have conflicting and asymmetric goals while complex cooperation within each team may exist. In addition, optimal solutions are impossible to compute using traditional dynamic programming approaches due to the intrinsic high dimensionality of the joint state space. Previously, in [6], where a team of defenders assumes that the attackers move towards their target in straight lines, a mixed-integer linear programming approach was used. [7] assumes that the attackers uses a linear feedback control law, and a mixed integer linear program was then relaxed into linear program. In complex pursuit-evasion games where players may change roles over time, a nonlinear model-predictive control [8] approach has been investigated. Approximate dynamic programming [9] has also been used to analyze reach-avoid games.

This work has been supported in part by NSF under CPS:ActionWebs (CNS-931843), by ONR under the HUNT (N0014-08-0696) and SMARTS (N00014-09-1-1051) MURIs and by grant N00014-12-1-0609, by AFOSR under the CHASE MURI (FA9550-10-1-0567).

M. Chen, and C. J. Tomlin are with the Department of Electrical Engineering and Computer Sciences, University of California, Berkeley, CA 94720, USA {mochen72, tomlin}@eecs.berkeley.edu

Z. Zhou is with the Department of Electrical Engineering, Stanford University, Stanford, CA 94305, USA zyzhou@stanford.edu

We thank Haomiao Huang for sharing MatLab code for 4D HJI calculations.

Although the above techniques provide some useful insight, they only work well when strong assumptions are made or when accurate models of the opposing team can be obtained. To solve general reach-avoid games, the Hamilton-Jacobi-Isaacs (HJI) approach [10] is ideal when the game is low-dimensional. The approach involves solving an HJI partial differential equation (PDE) in the joint state space of the players to compute a reach-avoid set, which partitions the joint state space of the players into a winning region for the defending team and one for the attacking team. The optimal strategies can then be extracted from the gradient of the solution. This approach is particularly useful because of the numerical tools [11], [12], [13] available, and has been able to solve a variety of problems [2], [11], [14]. The HJI approach can be applied to a large variety of player dynamics and does not explicitly assume any control strategy or prediction models for the players. However, the approach cannot be directly applied to our multiplayer reach-avoid game because its complexity scales exponentially with the number of players, making the approach only tractable for the two-player game. Thus, complexity-optimality trade-offs must be made.

For the two-player reach-avoid game, we first present the two-player HJI solution [2], which computes a four-dimensional (4D) reach-avoid set that determines which player wins the game assuming both players use the closed-loop optimal control strategy. Next, we propose the “path defense” approximation to the HJI solution, in which the defenders utilize a “semi-open-loop” control strategy. Here, we approximate two-dimensional (2D) slices of the reach-avoid sets by solving 2D Eikonal equations, and provide guarantees for the defending team’s performance.

For the multiplayer reach-avoid game, we propose to merge the N^2 pairwise outcomes using the graph theoretic maximum matching, which can be efficiently computed by known algorithms [15], [16]. When players on each team have identical dynamics, only a *single* HJI PDE needs to be solved to characterize *all* pairwise outcomes. Furthermore, when applying maximum matching to the two-player path defense solution, the computational complexity scales linearly with the number of players, as opposed to quadratically in the HJI approach.

II. THE REACH-AVOID PROBLEM

A. The Multiplayer Reach-Avoid Game

Consider $N_A + N_D$ players partitioned into the set of N_A attackers, $\{P_{A_i}\}_{i=1}^{N_A} = \{P_{A_1}, P_{A_2}, \dots, P_{A_{N_A}}\}$ and the set of N_D defenders, $\{P_{D_i}\}_{i=1}^{N_D} = \{P_{D_1}, \dots, P_{D_{N_D}}\}$, whose states are confined in a bounded, open domain $\Omega \subset \mathbb{R}^2$. The domain Ω is partitioned into $\Omega = \Omega_{\text{free}} \cup \Omega_{\text{obs}}$, where Ω_{free} is a compact set representing the free space in which the players can move, while $\Omega_{\text{obs}} = \Omega \setminus \Omega_{\text{free}}$ corresponds to obstacles in the domain.

Let $x_{A_i}, x_{D_j} \in \mathbb{R}^2$ denote the state of players P_{A_i} and P_{D_j} , respectively. Then given initial conditions $x_{A_i}^0 \in \Omega_{\text{free}}, i = 1, 2, \dots, N_A, x_{D_i}^0 \in \Omega_{\text{free}}, i = 1, 2, \dots, N_D$, we assume the dynamics of the players to be defined by the following decoupled system for $t \geq 0$:




Fig. 1: The components of a multiplayer reach-avoid game.

$$\begin{aligned} \dot{x}_{A_i}(t) &= v_{A_i} a_i(t), & x_{A_i}(0) &= x_{A_i}^0, i = 1, 2, \dots, N_A \\ \dot{x}_{D_i}(t) &= v_{D_i} d_i(t), & x_{D_i}(0) &= x_{D_i}^0, i = 1, 2, \dots, N_D \end{aligned} \quad (1)$$

where v_{A_i}, v_{D_i} denote maximum speeds for P_{A_i} and P_{D_i} respectively, and a_i, d_i denote controls of P_{A_i} and P_{D_i} respectively. We assume that a_i, d_i are drawn from the set $\Sigma = \{\sigma: [0, \infty) \rightarrow \bar{B}_n \mid \sigma \text{ is measurable}\}$, where \bar{B}_n denotes the closed unit disk in \mathbb{R}^2 . We also constrain the players to remain within Ω_{free} for all time. Denote the joint state of all players by $\mathbf{x} = (\mathbf{x}_A, \mathbf{x}_D)$ where $\mathbf{x}_A = (x_{A_1}, \dots, x_{A_{N_A}})$ is the attacker joint state $\{P_{A_i}\}_{i=1}^{N_A}$, and $\mathbf{x}_D = (x_{D_1}, \dots, x_{D_{N_D}})$ is the defender joint state $\{P_{D_i}\}_{i=1}^{N_D}$.

The attacking team wins whenever M of the N_A attackers reach some target set without being captured by the defenders; M is pre-specified with $0 < M \leq N_A$. The target set is denoted $\mathcal{T} \subset \Omega_{\text{free}}$ and is compact. The defending team wins if it can prevent the attacking team from winning by capturing or indefinitely delaying $N_A - M + 1$ attackers from reaching \mathcal{T} . An illustration of the game setup is shown in Fig. 1.

We will define the capture sets, \mathcal{C}_{ij} , to be $\mathcal{C}_{ij} = \{\mathbf{x} \in \Omega^{N_A+N_D} \mid \|x_{A_i} - x_{D_j}\|_2 \leq R_C\}$. P_{A_i} is captured by P_{D_j} if P_{A_i} 's position is within a distance R_C of P_{D_j} 's position.

In this paper, we address the following problems:

- 1) Given \mathbf{x}^0 , \mathcal{T} , and some fixed integer $M, 0 < M \leq N_A$, can the attacking team win?
- 2) More generally, given \mathbf{x}^0 and \mathcal{T} , how many attackers can the defending team prevent from reaching the target?

B. The Two-Player Reach-Avoid Game

We will answer the above questions about the general N_A vs. N_D reach-avoid game by using the solution to the two-player 1 vs. 1 game as a building block. In the two-player game, we denote the attacker as P_A , the defender as P_D , their states as x_A, x_D , and their initial conditions as x_A^0, x_D^0 . Their dynamics are

$$\begin{aligned} \dot{x}_A(t) &= v_A a(t), & x_A(0) &= x_A^0, \\ \dot{x}_D(t) &= v_D d(t), & x_D(0) &= x_D^0 \end{aligned} \quad (2)$$

The players' joint state becomes $\mathbf{x} = (x_A, x_D)$, and their joint initial condition becomes $\mathbf{x}^0 = (x_A^0, x_D^0)$. The capture set becomes simply $\mathcal{C} = \{(x_A, x_D) \in \Omega^2 \mid \|x_A - x_D\|_2 \leq R_C\}$.

P_A wins if it reaches the target \mathcal{T} without being captured by P_D . P_D wins if it can prevent P_A from winning by capturing

P_A or indefinitely delaying P_A from reaching \mathcal{T} . For the two-player reach-avoid game, we seek to answer the following:

- 1) Given \mathbf{x}^0 and \mathcal{T} , is the defender guaranteed to win?
- 2) More generally, given x_A and \mathcal{T} , what is the set of initial positions such that the defender is guaranteed to win?

III. THE HJI SOLUTION OF THE TWO-PLAYER GAME

The HJI approach for solving differential games is outlined in [2], [11], [17]. The optimal joint closed-loop control strategies for the attacker and the defender in a two-player reach-avoid game can be obtained by solving a 4D HJI PDE. This solution allows us to determine whether the defender will win against the attacker in a 1 vs. 1 setting.

In the two-player game, the attacker aims to reach \mathcal{T} while avoiding \mathcal{C} . Both players also avoid Ω_{obs} . In particular, the defender wins if the attacker is in Ω_{obs} , and vice versa. Therefore, we define the terminal set and avoid set to be

$$\begin{aligned} R &= \{\mathbf{x} \in \Omega^2 \mid x_A \in \mathcal{T}\} \cup \{\mathbf{x} \in \Omega^2 \mid x_D \in \Omega_{\text{obs}}\} \\ A &= \mathcal{C} \cup \{\mathbf{x} \in \Omega^2 \mid x_A \in \Omega_{\text{obs}}\} \end{aligned} \quad (3)$$

Given (3), we can define the corresponding implicit surface functions ϕ_R, ϕ_A required for solving the HJI PDE. Since $\Omega \subset \mathbb{R}^2$, the result is $\mathcal{RA}_{\infty}(R, A) \in \mathbb{R}^4$, a 4D reach-avoid set. If $\mathbf{x}^0 \in \mathcal{RA}_{\infty}(R, A)$, then the attacker is guaranteed to win the game by using the optimal control *even if* the defender is also using the optimal control; if $\mathbf{x}^0 \notin \mathcal{RA}_{\infty}(R, A)$, then the defender is guaranteed to win the game by using the optimal control *even if* the attacker is also using the optimal control.

IV. THE PATH DEFENSE SOLUTION TO THE TWO-PLAYER GAME

We approximate 2D slices of the 4D reach-avoid set (or simply "2D slices") in the path defense approach. Each slice will be taken at an attacker position. Here, we will assume that the defender is not slower than the attacker: $v_A \leq v_D$.

A. The Path Defense Game

The *Path Defense Game* is a two-player reach-avoid game in which the boundary of the target set is the shortest path between two points on $\partial\Omega$, and the target set is on one side of that shortest path (Fig. 2). We denote the target set as $\mathcal{T} = \Omega \setminus \mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$ for two given points on the boundary e_a, e_b . $\mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$ and $\bar{\mathcal{P}}(e_a, e_b)$ are defined below.

Definition 1: Path of defense. A path of defense, $\bar{\mathcal{P}}(e_a, e_b)$, is the shortest path between two boundary points $e_a, e_b \in \partial\Omega$. e_a and e_b are called the **anchor points** of path $\bar{\mathcal{P}}(e_a, e_b)$.

Denote the shortest path between *any* two points $x, y \in \Omega_{\text{free}}$ to be $\mathcal{P}(x, y)$, with length $\text{dist}(x, y)$, and requiring the attacker and defender $t_A(x, y), t_D(x, y)$ to traverse, respectively. We will also use $\text{dist}(\cdot, \cdot)$ with one or both arguments being sets in Ω to denote the shortest distance between the arguments.

Definition 2: Attacker's side of the path. A path of defense $\bar{\mathcal{P}}(e_a, e_b)$ partitions the domain Ω into two regions. Define $\mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$ to be the region that contains the attacker, not including points on the path $\bar{\mathcal{P}}(e_a, e_b)$. The attacker seeks to reach the target set $\mathcal{T} = \Omega \setminus \mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$.




Fig. 2: The components of a path defense game.

B. Solving The Path Defense Game

A path defense game can be directly solved by computing a 4D reach-avoid set. Since the direct solution is time- and memory- intensive, we propose an efficient approximation of 2D slices that is conservative towards the defender.

Definition 3: Strongly defendable path. $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable if regardless of the attacker's actions, the defender has a control function $d(\cdot)$ to reach $\bar{\mathcal{P}}(e_a, e_b)$ after finite time and prevent the attacker from reaching $\bar{\mathcal{P}}(e_a, e_b)$.

Definition 4: Attacker level set image. Given attacker position $x_A(t)$, define the attacker level set image with respect to anchor point e_a to be $x_A'(t; e_a) = \{x \in \bar{\mathcal{P}}(e_a, e_b) : t_A(x, e_a) = t_A(x_A(t), e_a)\}$. x_A' is the unique point on $\bar{\mathcal{P}}(e_a, e_b)$ such that $t_A(x_A', e_a) = t_A(x_A, e_a)$. Define $x_A'(t; e_b)$ similarly by replacing e_a with e_b . For convenience, we sometimes omit the time argument and write $x_A'(e_a)$.

Proposition 1: $\text{dist}(x_A'(e_b), e_a) \leq \text{dist}(x_A'(e_a), e_a)$.

Proof: First note that

$$\begin{aligned} \text{dist}(e_a, e_b) &\leq \text{dist}(x_A, e_a) + \text{dist}(x_A, e_b) \\ &= \text{dist}(x_A'(e_a), e_a) + \text{dist}(x_A'(e_b), e_b) \end{aligned}$$

Then, since the left hand side is given by $\text{dist}(e_a, e_b) = \text{dist}(e_a, x_A'(e_b)) + \text{dist}(x_A'(e_b), e_b)$, the result follows. ■

Definition 5: Capture set: Define the capture set to be $D_C(y, t) = \{x \mid \|x - y(t)\|_2 \leq R_C\}$. We will drop the second argument of D_C when y does not depend on time.

Remark 1: Given $\bar{\mathcal{P}}(e_a, e_b)$, suppose the attacker level set image is within defender's capture set at some time s :

$$x_A'(s; e_a) \in D_C(x_D, s) \quad (\text{or} \quad x_A'(s; e_b) \in D_C(x_D, s))$$

Then, there exists a control for the defender to keep the attacker level set image within the capture radius of the defender thereafter:

$$\begin{aligned} x_A'(t; e_a) &\in D_C(x_D, t) \forall t \geq s \\ (\text{or } x_A'(t; e_b) &\in D_C(x_D, t) \forall t \geq s) \end{aligned}$$

This is because the attacker level set image can move at most as fast as the attacker, who is not faster than the defender.

Definition 6: Regions induced by point p on path. Given a point $p \in \bar{\mathcal{P}}(e_a, e_b)$, define a region $\mathcal{R}_a(p)$ associated the point p and anchor point e_a as follows:

$$\mathcal{R}_a(p) = \{x : \text{dist}(x, e_a) \leq \text{dist}(D_C(p), e_a)\} \quad (4)$$

Define $\mathcal{R}_b(p)$ similarly by replacing e_a with e_b .

Lemma 1: Suppose $x_D^0 = p \in \bar{\mathcal{P}}(e_a, e_b)$, and $v_A = v_D$. Then, $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable if and only if x_A^0 is outside the region induced by p : $x_A^0 \in \Omega \setminus (\mathcal{R}_a \cup \mathcal{R}_b)$.




Fig. 3: Left: If the attacker is in $\mathcal{R}_a \cup \mathcal{R}_b$ and moves towards e_a , he will be able to reach e_a without being captured. Right: If the attacker is not in $\mathcal{R}_a \cup \mathcal{R}_b$, there is no point on the path $p' \in \bar{\mathcal{P}}(e_a, e_b)$ that he can reach without being captured.

Proof: See Fig. 3. Assume $x_A^0 \notin \mathcal{T} = \Omega \setminus \mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$, otherwise the attacker would start inside the target set.

First, we show that if $x_A^0 \in \mathcal{R}_a \cup \mathcal{R}_b$, then the attacker can reach e_a or e_b and hence $\Omega \setminus \mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$ without being captured. Without loss of generality (WLOG), suppose $x_A^0 \in \mathcal{R}_a$. To capture the attacker, the defender's capture set must contain $x_A'(e_a)$ or $x_A'(e_b)$ at some time t . By Definition 6, we have $\text{dist}(x_A'(e_a), e_a) < \text{dist}(D_C(p), e_a)$, so $t_A(x_A'(e_a), e_a) < t_D(D_C(p), e_a)$. By Proposition 1, $\text{dist}(x_A'(e_b), e_a) \leq \text{dist}(x_A'(e_a), e_a)$, so it suffices to show that the defender's capture set cannot reach $x_A'(e_a)$ before the attacker reaches e_a .

If the attacker moves towards e_a along $\mathcal{P}(x_A^0, e_a)$ with maximum speed, then $x_A'(e_a)$ will move towards e_a along $\mathcal{P}(x_A'(e_a), e_a)$ at the same speed. Since $t_A(x_A, e_a) = t_A(x_A'(e_a), e_a) < t_D(D_C(p), e_a)$, x_A will reach e_a before the defender capture set $D_C(x_D, t)$ does.

Next we show, by contradiction, that if $x_A \notin \mathcal{R}_a \cup \mathcal{R}_b$, then the attacker cannot reach $\Omega \setminus \mathcal{S}_A(\bar{\mathcal{P}}(e_a, e_b))$ without being captured. Suppose P_A will reach some point p' before $D_C(P_D, t)$ does, i.e. $\text{dist}(x_A^0, p') < \text{dist}(D_C(x_D^0), p') = \text{dist}(D_C(p), p')$. WLOG, assume $p' \in \mathcal{P}(p, e_b)$, and note that $\text{dist}(D_C(p), e_b) < \text{dist}(x_A^0, e_b)$ since the attacker is not in \mathcal{R}_b . Starting with the definition of the shortest path, we have

$$\begin{aligned} \text{dist}(x_A^0, e_b) &\leq \text{dist}(x_A^0, p') + \text{dist}(p', e_b) \\ &< \text{dist}(D_C(p), p') + \text{dist}(p', e_b) \\ &= \text{dist}(D_C(p), e_b) \\ \text{dist}(x_A^0, e_b) &< \text{dist}(x_A^0, e_b) \quad (\text{since } x_A^0 \notin \mathcal{R}_a) \end{aligned} \quad (5)$$

This is a contradiction. Therefore, the attacker cannot cross any point p' on $\bar{\mathcal{P}}(e_a, e_b)$ without being captured. ■

If $v_A < v_D$, P_A being outside of $\mathcal{R}_a \cup \mathcal{R}_b$ becomes a sufficient condition for the strong defendability of $\bar{\mathcal{P}}(e_a, e_b)$.

In general, x_D^0 may not be on $\bar{\mathcal{P}}(e_a, e_b)$. In this case, if the defender can arrive at p before the attacker moves into $\mathcal{R}_a(p) \cup \mathcal{R}_b(p)$, then $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable. Thus, given $x^0 = (x_A^0, x_D^0)$, we may naively check whether a path $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable by checking whether there exists some $p \in \bar{\mathcal{P}}(e_a, e_b)$ such that $t_D(x_D^0, p) \leq t_A(x_A^0, \mathcal{R}_a(p) \cup \mathcal{R}_b(p))$. If so, then $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable. The next lemma shows that it is necessary and sufficient to check whether *one* special point, $p^* \in \bar{\mathcal{P}}(e_a, e_b)$, can be the first arrival point for strongly defending $\bar{\mathcal{P}}(e_a, e_b)$.

Remark 2: Given $p \in \bar{\mathcal{P}}(e_a, e_b)$, $\text{dist}(x_A^0, \mathcal{R}_a(p)) = \text{dist}(x_A^0, e_a) - \text{dist}(D_C(p), e_a)$. Similarly, $\text{dist}(x_A^0, \mathcal{R}_b(p)) = \text{dist}(x_A^0, e_b) - \text{dist}(D_C(p), e_b)$.

Lemma 2: Define $p^* \in \bar{\mathcal{P}}(e_a, e_b)$ such that $t_A(x_A^0, \mathcal{R}_a) = t_A(x_A^0, \mathcal{R}_b)$. Then, $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable if and only if the defender can defend $\bar{\mathcal{P}}(e_a, e_b)$ by first going to p^* .

Proof: One direction is clear by definition.

We now show the other direction's contrapositive: if the defender cannot defend $\bar{\mathcal{P}}(e_a, e_b)$ by first going to p^* , then $\bar{\mathcal{P}}(e_a, e_b)$ is not strongly defendable. Equivalently, we show that if choosing p^* as the first entry point does not allow the defender to defend $\bar{\mathcal{P}}(e_a, e_b)$, then no other entry point does.

Suppose that the defender cannot defend $\bar{\mathcal{P}}(e_a, e_b)$ by choosing p^* as the first entry point, but can defend $\bar{\mathcal{P}}(e_a, e_b)$ by choosing another entry point p' . WLOG, assume $\text{dist}(D_C(p^*), e_a) > \text{dist}(D_C(p'), e_a)$. This assumption moves p' further away from e_a than p^* , causing \mathcal{R}_a to move closer to x_A^0 . Starting with Remark 2, we have

$$\begin{aligned} \text{dist}(x_A^0, \mathcal{R}_a(p^*)) &= \text{dist}(x_A^0, e_a) - \text{dist}(D_C(p^*), e_a) \\ t_A(x_A^0, \mathcal{R}_a(p^*)) &= t_A(x_A^0, e_a) - t_A(D_C(p^*), e_a) \end{aligned} \quad (6)$$

Similarly, for the point p' , we have

$$\begin{aligned} \text{dist}(x_A^0, \mathcal{R}_a(p')) &= \text{dist}(x_A^0, e_a) - \text{dist}(D_C(p'), e_a) \\ t_A(x_A^0, \mathcal{R}_a(p')) &= t_A(x_A^0, e_a) - t_A(D_C(p'), e_a) \end{aligned} \quad (7)$$

Then, subtracting the above two equations, we see that the attacker can get to \mathcal{R}_a sooner by the following amount:

$$\begin{aligned} t_A(x_A^0, \mathcal{R}_a(p^*)) - t_A(x_A^0, \mathcal{R}_a(p')) &= t_A(D_C(p'), e_a) - t_A(D_C(p^*), e_a) \\ &= t_A(p', p^*) \geq t_D(p', p^*) \end{aligned} \quad (8)$$

We now show that the defender can get to p' sooner than to p^* by less than the amount $t_D(p', p^*)$, and in effect “gains less time” than the attacker does by going to p' . We assume that p' is closer to the defender than p^* is (otherwise the defender “loses time” by going to p'). By the triangle inequality,

$$\begin{aligned} \text{dist}(x_D^0, p^*) &\leq \text{dist}(x_D^0, p') + \text{dist}(p', p^*) \\ \text{dist}(x_D^0, p^*) - \text{dist}(x_D^0, p') &\leq \text{dist}(p', p^*) \\ t_D(x_D^0, p^*) - t_D(x_D^0, p') &\leq t_D(p', p^*) \end{aligned} \quad (9)$$

Lemmas 1 and 2 give a simple algorithm to compute, given x_A^0 , the region that the defender must be in for a path of defense $\bar{\mathcal{P}}(e_a, e_b)$ to be strongly defendable:

- 1) Given e_a, e_b, x_A^0 , compute p^* and $\mathcal{R}_a(p^*), \mathcal{R}_b(p^*)$.
- 2) If $v_A = v_D$, then $\bar{\mathcal{P}}(e_a, e_b)$ is strongly defendable if and only if $x_D^0 \in \mathcal{D}(e_a, e_b; x_A^0) = \{x : t_D(x, p^*) \leq t_A(x_A^0, \mathcal{R}_a \cup \mathcal{R}_b)\}$.

The computations in this algorithm can be efficiently done by solving a series of 2D Eikonal equations by using FMM [12], reducing our 4D problem to 2D. Fig. 4 illustrates the proof of Lemma 2 and the defender winning region \mathcal{D} .




Fig. 4: Left: If the defender cannot defend $\bar{\mathcal{P}}(e_a, e_b)$ by first going to p^* , then he cannot defend $\bar{\mathcal{P}}(e_a, e_b)$ by going to any other point p . Right: Defender winning region \mathcal{D} .

C. The Path Defense Solution to the Reach-Avoid Game

The central idea in using path defense is that if the target set is enclosed by some strongly defendable path for some e_a, e_b , then the defender can win the game using the semi-open-loop strategy outlined in this section, *even if* the attacker uses the optimal control. Checking for strongly defendable paths adds more conservatism towards the defender, but makes computation much more efficient.

Naively, one could fix e_a , then search all other anchor points $e_b \in \partial\Omega$ to find a defendable path. However, we can reduce the number of paths that needs to be checked by only checking paths of defense $\bar{\mathcal{P}}(e_a, e_b)$ that touch the target set. In a simply connected domain, this reduction in the number of paths checked does not introduce any additional conservatism.

If some strongly defendable path $\bar{\mathcal{P}}(e_a, e_b)$ encloses the target set, then the defender's strategy would be to first go to $p^* \in \bar{\mathcal{P}}(e_a, e_b)$ (an open-loop strategy), then move towards $x_{A'}(e_a)$ or $x_{A'}(e_b)$ until the level set image is captured (a closed-loop strategy). Finally, the defender can simply track the captured level set image (a closed-loop strategy). This is a “semi-open-loop” strategy. The following algorithm approximates a 2D slice conservatively towards the defender:

Algorithm 1: Given attacker position,

- 1) Choose some point $e_a \in \partial\Omega$, which defines e_b to create a path of defense $\bar{\mathcal{P}}(e_a, e_b)$ that touches the target \mathcal{T} .
- 2) Repeat step 1 for a desired set of points $e_a \in \partial\Omega$.
- 3) For some particular $\bar{\mathcal{P}}(e_a, e_b)$, determine the defender winning region $\mathcal{D}(e_a, e_b; x_A^0)$.
- 4) Repeat step 3 for all the paths created in steps 1 and 2.
- 5) The union $\bigcup_{e_a} \mathcal{D}(e_a, e_b; x_A^0)$ gives the approximate 2D slice, representing the conservative winning region for the defender in the two-player reach-avoid game.

V. FROM TWO-PLAYER TO MULTIPLAYER

A. Maximum Matching

We piece together the outcomes of all attacker-defender pairs using maximum matching as follows:

Algorithm 2:

- 1) Construct a bipartite graph with two sets of nodes $\{P_{A_i}\}_{i=1}^{N_A}, \{P_{D_i}\}_{i=1}^{N_D}$. Each node represents a player.
- 2) For all i, j , draw an edge between P_{D_i} and P_{A_j} if P_{D_i} wins against P_{A_j} in a two-player reach-avoid game.
- 3) Run any matching algorithm (eg. [15], [16]) to find a maximum matching in the graph.




Fig. 5: An illustration of using maximum matching to conservatively approximate the multiplayer reach-avoid game.

After finding a maximum matching, we can guarantee an upper bound on the number of attackers that will be able to reach the target. If the maximum matching is of size m , then the defending team would be able to prevent *at least* m attackers from reaching the target.

B. Time-Varying Defender-Attacker Pairings

With the next algorithm, the bipartite graph and its corresponding maximum matching can be updated, potentially in real time, as the players change positions during the game:

Algorithm 3:

- 1) Given each x_{D_i} and each x_{A_j} , determine whether P_{A_j} can win against P_{D_i} for all i, j .
- 2) Assign a defender to each attacker that is part of a maximum matching.
- 3) For a short duration Δ , apply a winning control input and compute the resulting trajectory for each defender that is part of the maximum matching. For the rest of the defenders and for all attackers, compute the trajectories assuming some (any) control function.
- 4) Update the player positions after the duration Δ and repeat steps 1 to 3 with the new player positions.

As $\Delta \rightarrow 0$, the above procedure continuously computes a bipartite graph and its maximum matching. As long as each defender uses a winning control input against the paired-up attacker, the size of maximum matching can never decrease.

C. Application to the Two-Player HJI Solution

In general, solving $N_A N_D$ 4D HJI PDEs gives us the pairwise outcomes between every attacker-defender pair. The computation time required is thus $C N_A N_D$, where C is the time required to solve a single 4D HJI PDE. The pairwise outcomes can then be merged together to approximate the N_A vs. N_D game. In the case where each team has a single maximum speed, solving *one* 4D HJI PDE would characterize all pairwise outcomes.

Since the solution to the 4D HJI PDE characterizes pairwise outcomes based on any attacker-defender joint-state, it allows for real-time updates of the maximum matching. As players move to new positions, the pairwise outcome can be updated by simply checking whether (x_{A_i}, x_{D_j}) is in $\mathcal{RA}_\infty(R, A)$.

D. Application to the Two-Player Path Defense Solution

To use the pairwise outcomes determined by the path defense approach for approximating the solution to the multiplayer game, we add the following step to Algorithm 1:

- 6) Repeat steps 3 to 5 for every attacker position.

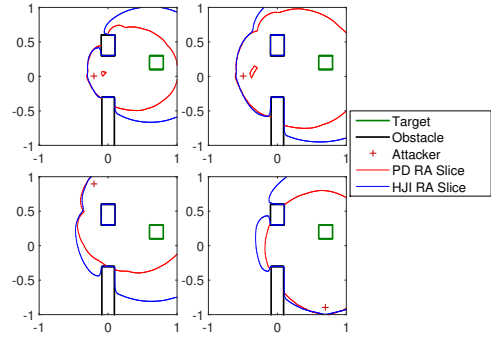


Fig. 6: Reach-avoid slices computed using the HJI approach and the path defense approach.

For a given domain, set of obstacles, and target set, steps 1 and 2 in Algorithm 1 only need to be performed once, regardless of the number of players. In step 3, the speeds of defenders come in only through a single distance calculation from p^* , which only needs to be done once per attacker position. Therefore, the total computation time required is on the order of $C_1 + C_2 N_A$, where C_1 is the time required for steps 1 and 2, C_2 is the time required for steps 3 to 5.

E. Defender Cooperation

VI. NUMERICAL RESULTS

We use a 4 vs. 4 example to illustrate our methods. The game is played on a square domain with obstacles. Defenders have a capture radius of 0.1 units, and all players have the same maximum speed. Computations were done on a laptop with a Core i7-2640M processor with 4 gigabytes of memory.

A. HJI Formulation

Fig. 6 shows the results of solving the corresponding 4D HJI PDE in blue. Computing the 4D reach-avoid set on a grid with 45 points in each dimension took approximately 30 minutes. All players have the same maximum speed, so only a single 4D HJI PDE needed to be solved. To visualize the 4D reach-avoid set, we take 2D slices of the 4D reach-avoid set sliced at each attacker's position.

Fig. 7 shows the bipartite graph and maximum matching obtained from the pairwise outcomes. The maximum matching is of size 4. This guarantees that if each defender plays optimally against the attacker matched by the maximum matching, then *no* attacker will be able to reach the target.

B. Path Defense Formulation

Fig. 6 shows, in red, the results of using the path defense approach to compute conservative approximations of the 4D reach-avoid set sliced at various attacker positions.

Fig. 7 shows the bipartite graph and maximum matching resulting from the pairwise outcomes. In this case, the maximum matching is of size 3. This guarantees that if each defender plays against the attacker matched by the maximum matching using the semi-open-loop strategy, then *at most* 1 attacker will be able to reach the target.




Fig. 7: Merging pairwise outcomes with maximum matching.




Fig. 8: Performance of the Path Defense Solution.

Computations were done on a 200×200 grid, and 937 paths were used to compute the results in Fig. 6. Computation time varies with the number of paths we chose in steps 1 and 2 in Algorithm 1. Taking the union of the defender winning regions from more paths will give a less conservative result, but requires more computation time. A summary of the performance is shown in Fig. 8. With 937 paths, the computation of paths took approximately 60 seconds, and the computation of the 2D slice given the set of paths took approximately 30 seconds. However, very few paths are needed to approximate a 2D slice: Even with as few as 30 paths, the computed 2D slice covers more than 95% of the area of the 2D slice computed using 937 paths. This reduces the computation time of the paths to 2.5 seconds, and the computation time of the 2D slices given the paths to 2.1 seconds.

C. Defender Cooperation

D. Real-Time Maximum Matching Updates

After determining all pairwise outcomes, pairwise outcomes of any joint state of the attacker-defender pair are characterized by the HJI approach. This allows for updates of the bipartite graph and its maximum matching as the players play out the game in real time. Fig. 9 shows the maximum matching at several time snapshots of a 4 vs. 4 game. Each defender that is part of a maximum matching plays optimally against the paired-up attacker, and the remaining defender plays optimally against the closest attacker. The attackers' strategy is to move towards the target along the shortest path while steering clear of the obstacles by 0.125 units. The maximum matching is updated every $\Delta = 0.005$ seconds. At $t = 0$ and $t = 0.2$, the maximum matching is of size 3, which guarantees that at most one attacker will be able to reach the target. After $t = 0.4$, the maximum matching size increases to 4, which guarantees that no attacker will be able to reach the target.

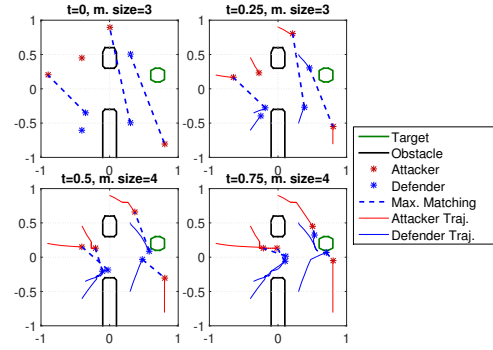


Fig. 9: Increasing maximum matching size over time.

VII. CONCLUSION AND FUTURE WORK

A multiplayer reach-avoid game is numerically intractable to analyze by directly solving the corresponding high dimensional HJI PDE. To address this, we presented a way to tie together pairwise outcomes using maximum matching to approximate the solution to the full multiplayer game, guaranteeing an upper bound on the number of attackers that can reach the target. We also presented two approaches for determining the pairwise outcomes. The HJI approach is computationally more expensive, produces the optimal closed-loop control strategy for each attacker-defender pair, and efficiently allows for real time maximum updates. The path defense approach is conservative towards the defender, performs computation on the state space of a single player as opposed to the joint state space, and scales only linearly in the number of players.

REFERENCES

- [1] H. Huang, "Reachability-based control for human and autonomous agents in adversarial games," Ph.D. dissertation, Stanford University, 2012.
- [2] H. Huang, J. Ding, W. Zhang, and C. Tomlin, "A differential game approach to planning in adversarial scenarios: A case study on capture-the-flag," in *Robotics and Automation (ICRA), 2011 IEEE International Conference on*, 2011, pp. 1451–1456.
- [3] Department of the Air Force, "United States Air Force unmanned aircraft systems flight plan 2009–2047," *oai.dtic.mil*, Jan 2009.
- [4] H. Erzberger, "Automated conflict resolution for air traffic control," *25th International Congress of the Aeronautical Sciences*, Jan 2006.
- [5] A. Madrigal, "Autonomous robots invade retail warehouses," <http://www.wired.com/wiredscience/2009/01/retailrobots/>.
- [6] M. Earl and R. D'Andrea, "A decomposition approach to multi-vehicle cooperative control," *Robotics and Autonomous Systems*, vol. 55, no. 4, pp. 276–291, 2007.
- [7] G. Chasparis and J. Shamma, "Linear-programming-based multi-vehicle path planning with adversaries," *Proceedings of the American Control Conference*, pp. 1072–1077, 2005.
- [8] J. Sprinkle, J. Eklund, H. Kim, and S. Sastry, "Encoding aerial pursuit/evasion games with fixed wing aircraft into a nonlinear model predictive tracking controller," *IEEE Conference on Decision and Control*, 2004.
- [9] J. McGrew, J. How, L. Bush, B. Williams, and N. Roy, "Air combat strategy using approximate dynamic programming," *AIAA Guidance, Navigation, and Control Conference*, Aug 2008.
- [10] R. Isaacs, *Differential Games*. New York: Wiley, 1967.
- [11] I. Mitchell, A. Bayen, and C. Tomlin, "A time-dependent Hamilton-Jacobi formulation of reachable sets for continuous dynamic games," *IEEE Transactions on Automatic Control*, vol. 50, no. 7, pp. 947–957, July 2005.

- [12] J. A. Sethian, "A fast marching level set method for monotonically advancing fronts," *Proceedings of the National Academy of Sciences*, vol. 93, no. 4, pp. 1591–1595, 1996. [Online]. Available: <http://www.pnas.org/content/93/4/1591.abstract>
- [13] S. Osher and R. Fedkiw, *Level Set Methods and Dynamic Implicit Surfaces*. Springer-Verlag, 2002, ISBN: 978-0-387-95482-0.
- [14] J. Ding, J. Sprinkle, S. S. Sastry, and C. J. Tomlin, "Reachability calculations for automated aerial refueling," in *IEEE Conference on Decision and Control*, Cancun, Mexico, 2008.
- [15] A. Schrijver, *Combinatorial Optimization : Polyhedra and Efficiency (Algorithms and Combinatorics)*. Springer, Jul. 2004. [Online]. Available: <http://www.worldcat.org/isbn/3540204563>
- [16] M. Karpinski and W. Rytter, *Fast parallel algorithms for graph matching problems*. New York, NY, USA: Oxford University Press, Inc., 1998.
- [17] I. Mitchell, *A Toolbox of Level Set Methods*, 2009, <http://people.cs.ubc.ca/~mitchell/ToolboxLS/index.html>.